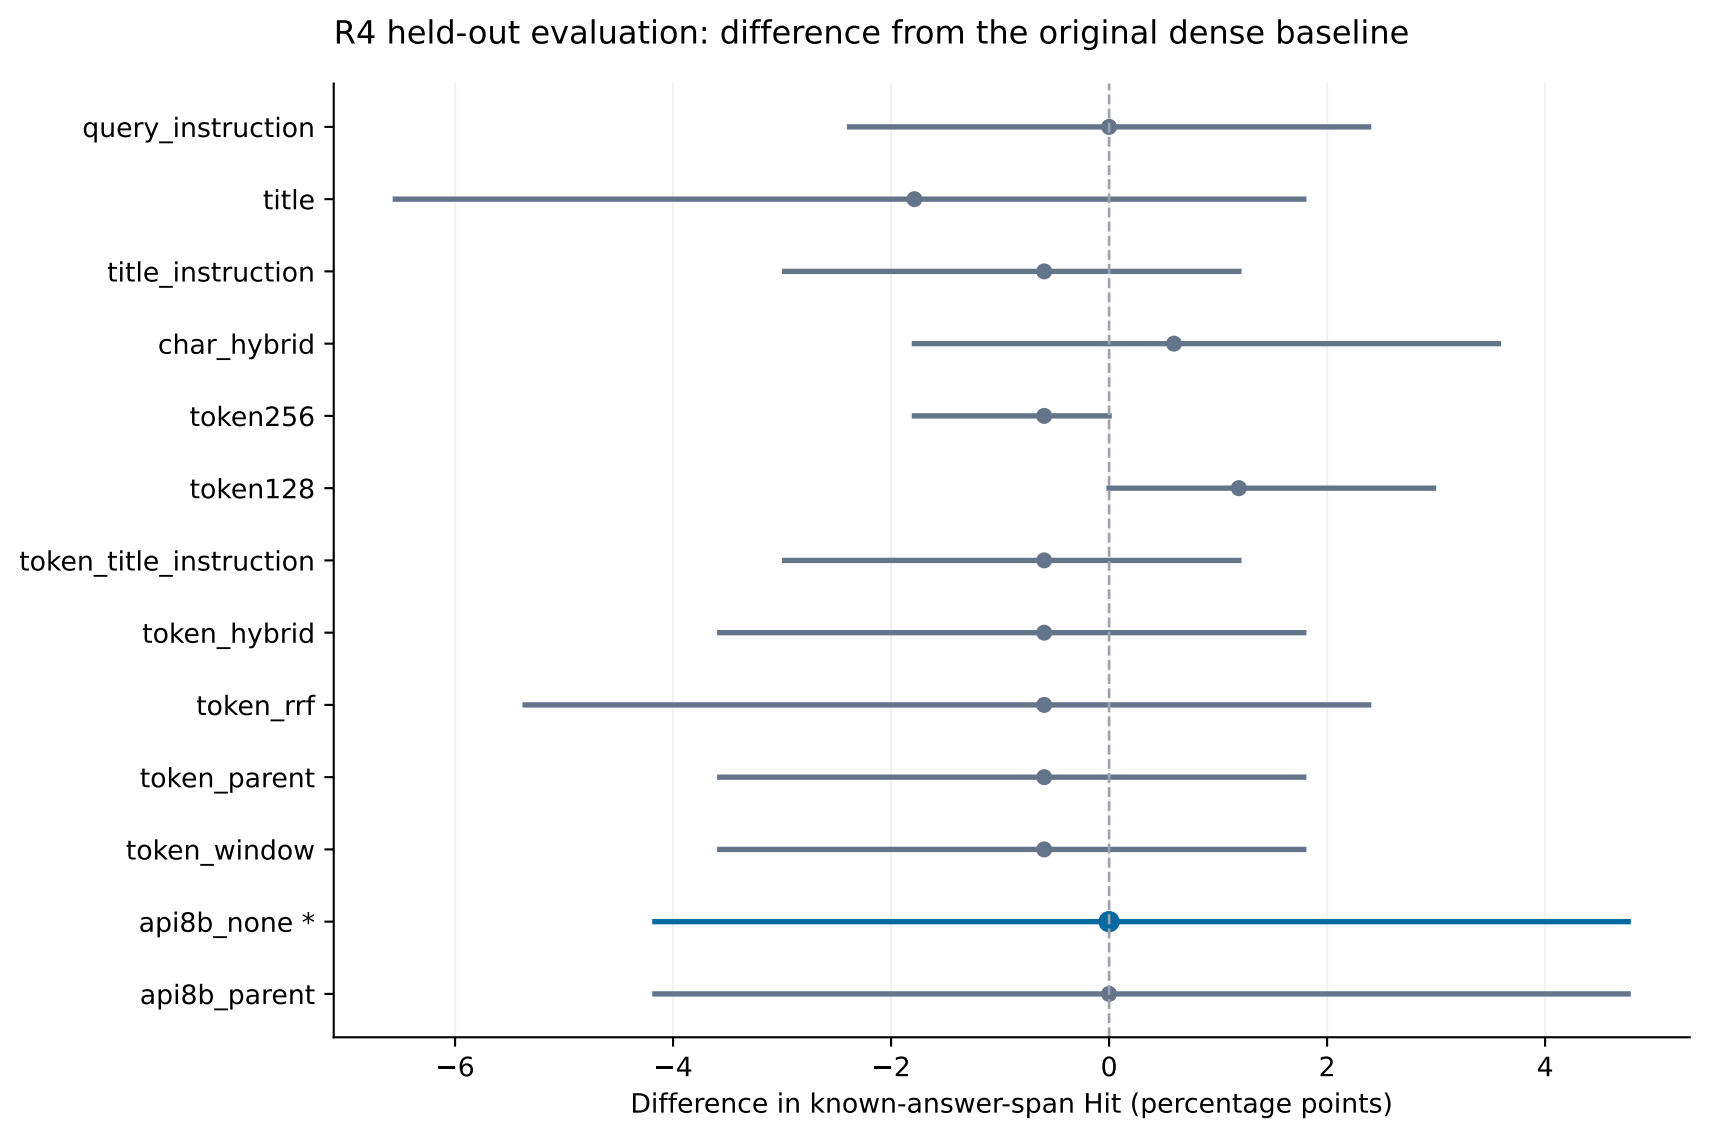

84 queries / 42 bilingual families; mixed + forced cross-language retrieval.
Budget: 1,024 Qwen tokenizer tokens of source text. Bars: paired family bootstrap 95% intervals.
* Selected on dev before held-out evaluation. Exploratory intervals; no multiple-comparison correction.
Source: public MLQA CS-related Wikipedia slice; not course-specific evidence annotations.