Fixed-rewrite retrieval fusion: ranking gains and regressions

R8A offline exploration · reused R7 test partition · not a new independent holdout

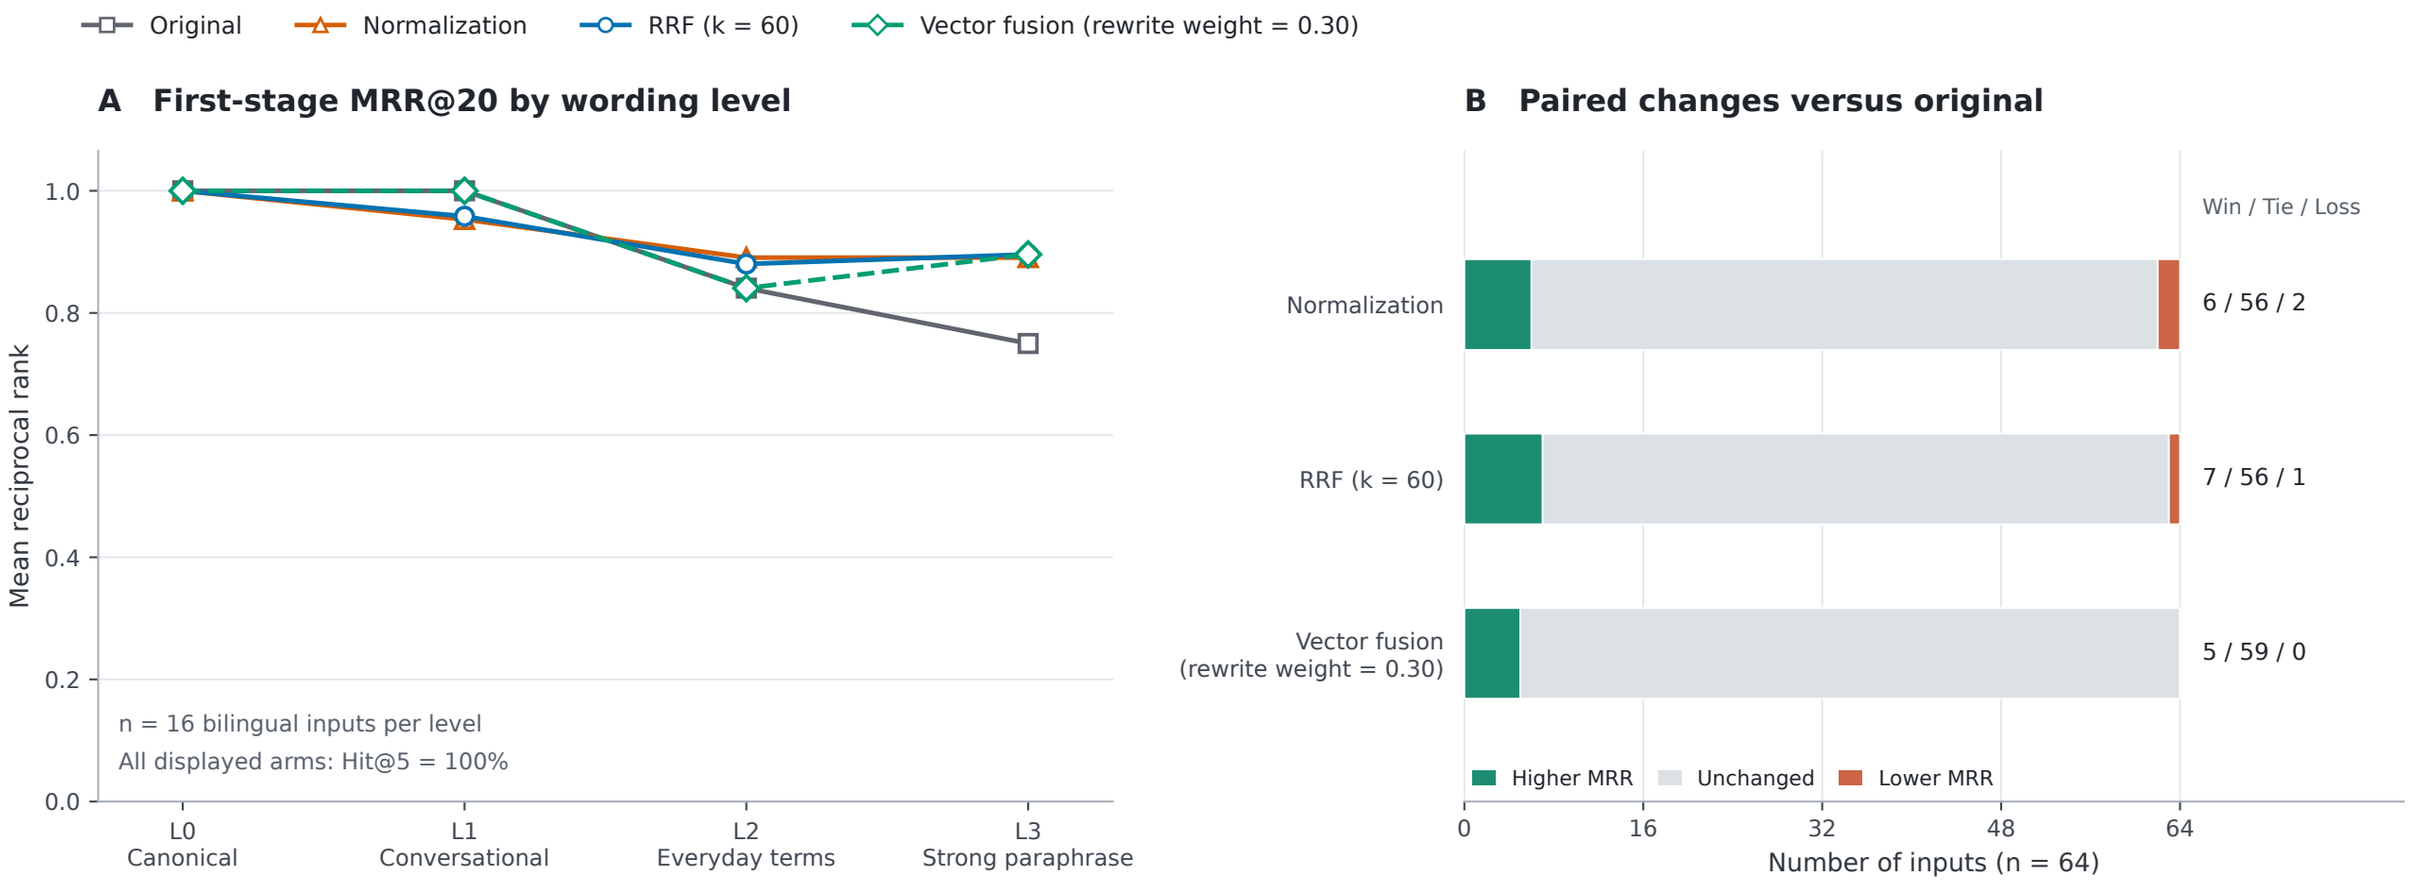

Same archived rewrite text, query embeddings, 380 chunks, and hybrid dense–lexical retrieval; at most 20 output candidates per arm.
Eight question families; AI-authored bilingual variants and AI-reviewed labels. L4 ambiguous inputs are excluded from retrieval scoring.
Descriptive point estimates only. The reused partition is exploratory, not fresh validation or a basis for choosing a production winner.
First-stage retrieval only: no reranking, answer generation, or latency comparison. New model/API calls: 0; new API cost: USD 0.
α = 0.30 mixes 70% original and 30% rewrite unit vectors; its lexical query remains original. Complete seven-arm results are in the report.