Selective clarification and first-stage retrieval

R8B · held-out questions on previously seen textbooks · simulated user replies

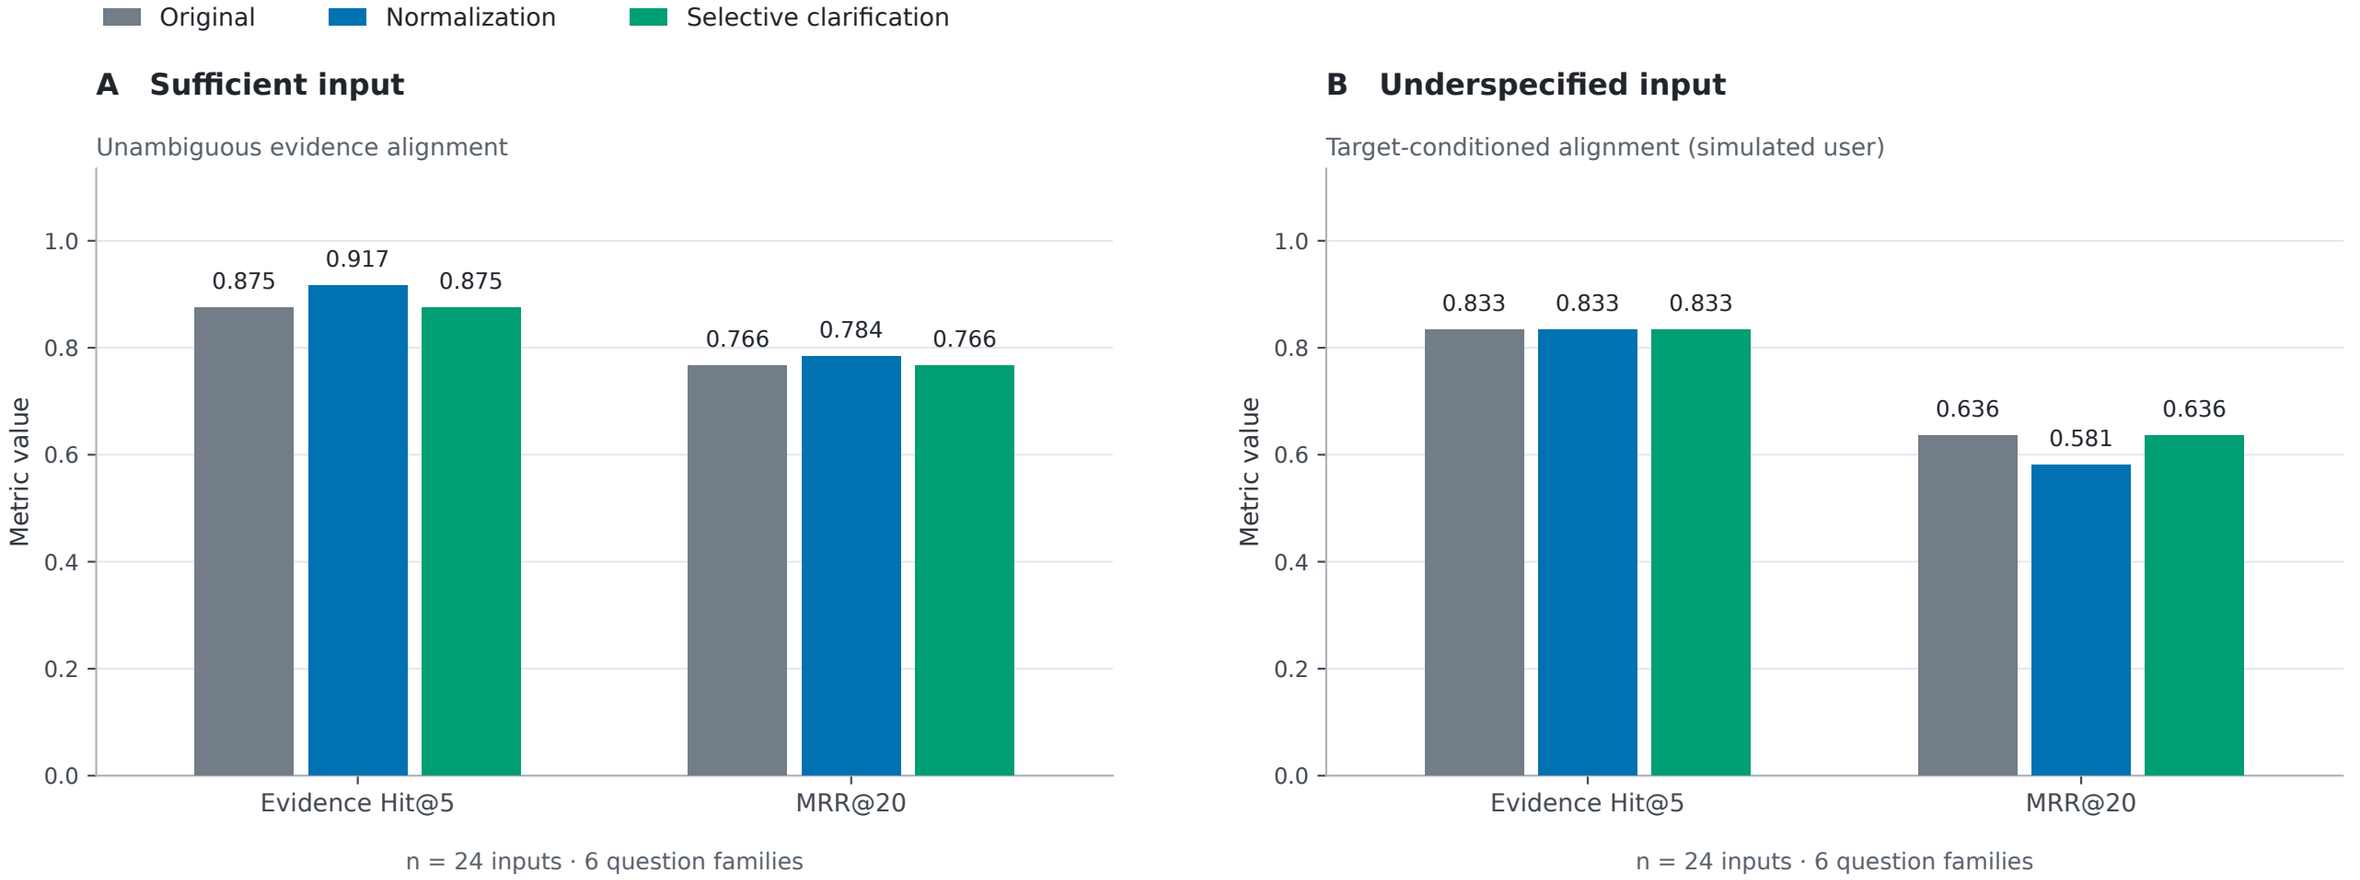

Selective policy: clarification decisions

Sufficient inputs: asked 0/24; resolved 0; unknown 0; unresolved asks 0.

Underspecified inputs: asked 1/24; resolved 1; unknown 0; unresolved asks 0.

48 held-out inputs in total; six families contribute to both groups. Languages and formal/colloquial styles remain in each denominator.
AI-authored inputs, AI-reviewed evidence labels, and simulated replies; this is not a real-student or teacher-gold evaluation.
Resolved is model-reported; unknown replies are a subset of unresolved asks. All unresolved, fallback, and failure cases are retained.
Panels use different interpretations and must not be pooled as a single accuracy score. Point estimates; paired intervals are in the report.
First-stage retrieval only, at most 20 candidates per policy. No reranking, answer generation, or end-to-end outcome measurement.